

DLP Lab6 Report

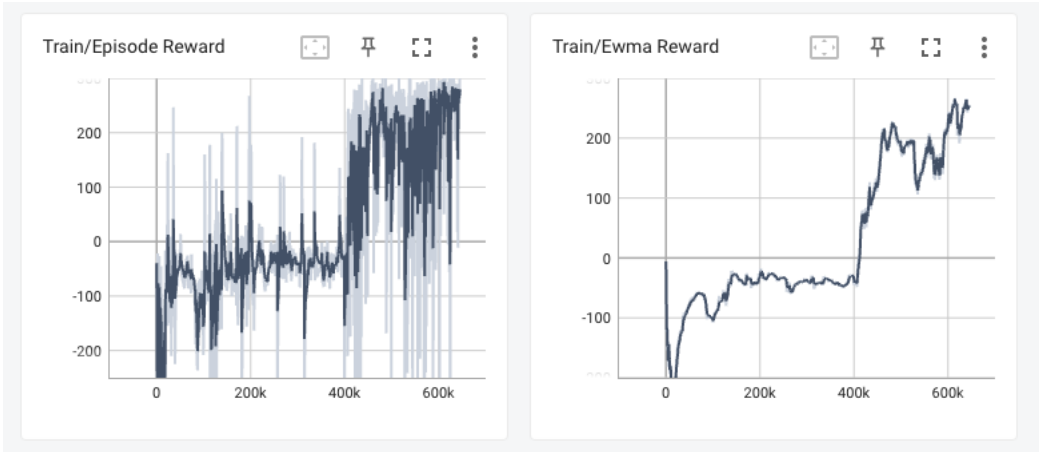
▼ 資科工碩 穆冠堯

Experiment Result

- LunarLander-v2
 - Testing Result

```
(dqn) minnie-mu@user:/eva_data6/minnie/NYCU/DLP/Lab6$ python dqn.py --test_only
/home/minnie-mu/anaconda3/envs/dqn/lib/python3.8/site-packages/gym/logger.py:30: UserWarning:
  warnings.warn(colorize('%s: %s'%( 'WARN', msg % args), 'yellow'))
Start Testing
Episode: 0      \Length: 176      Total Reward: 254.47
Episode: 1      \Length: 170      Total Reward: 289.52
Episode: 2      \Length: 190      Total Reward: 282.08
Episode: 3      \Length: 196      Total Reward: 284.03
Episode: 4      \Length: 276      Total Reward: 274.22
Episode: 5      \Length: 191      Total Reward: 267.53
Episode: 6      \Length: 157      Total Reward: 44.53
Episode: 7      \Length: 300      Total Reward: 267.99
Episode: 8      \Length: 221      Total Reward: 309.46
Episode: 9      \Length: 490      Total Reward: 262.55
Average Reward 253.63794992233366
```

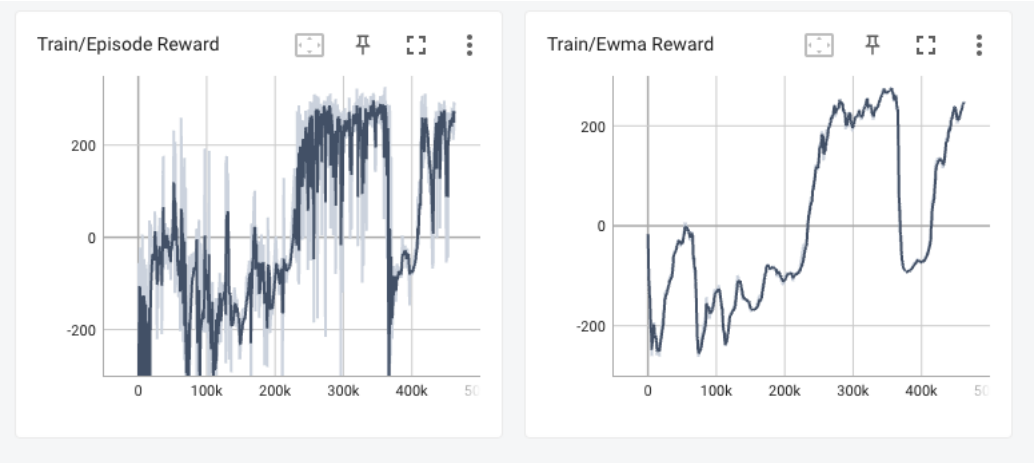
- Tensorboard



- LunarLander-Continuous
 - Testing Result

```
(dqn) minnie-mu@user:/eva_data6/minnie/NYCU/DLP/Lab6$ python ddpg.py --test_only
/home/minnie-mu/anaconda3/envs/dqn/lib/python3.8/site-packages/gym/logger.py:30:
casting to float32
  warnings.warn(colorize('%s: %s'%( 'WARN', msg % args), 'yellow'))
Start Testing
Episode: 0      Length: 194      Total reward: 229.71
Episode: 1      Length: 197      Total reward: 279.70
Episode: 2      Length: 827      Total reward: 224.76
Episode: 3      Length: 1000     Total reward: 138.31
Episode: 4      Length: 209      Total reward: 308.36
Episode: 5      Length: 934      Total reward: 232.69
Episode: 6      Length: 251      Total reward: 300.87
Episode: 7      Length: 200      Total reward: 295.66
Episode: 8      Length: 254      Total reward: 311.08
Episode: 9      Length: 1000     Total reward: 151.54
Average Reward 247.2662122894003
```

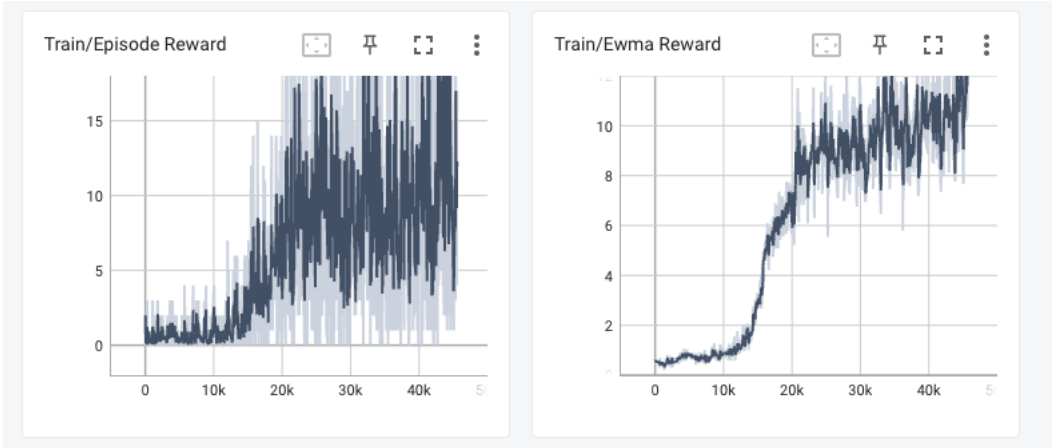
- Tensorboard



- BreakoutNoFrameskip-v4
 - Testing Result

```
(dqn) minnie-mu@user:/eva_data6/minnie/NYCU/DLP/Lab6$ python dqn_breakout_example.py --test_only
Start Testing
dqn_breakout_example.py:121: UserWarning: To copy construct from a tensor, it is recommended to use
, rather than torch.tensor(sourceTensor).
  state = torch.tensor(state, device=self.device)
episode 1: 361.00
episode 2: 433.00
episode 3: 428.00
episode 4: 439.00
episode 5: 409.00
episode 6: 428.00
episode 7: 428.00
episode 8: 394.00
episode 9: 453.00
episode 10: 449.00
Average Reward: 422.20
```

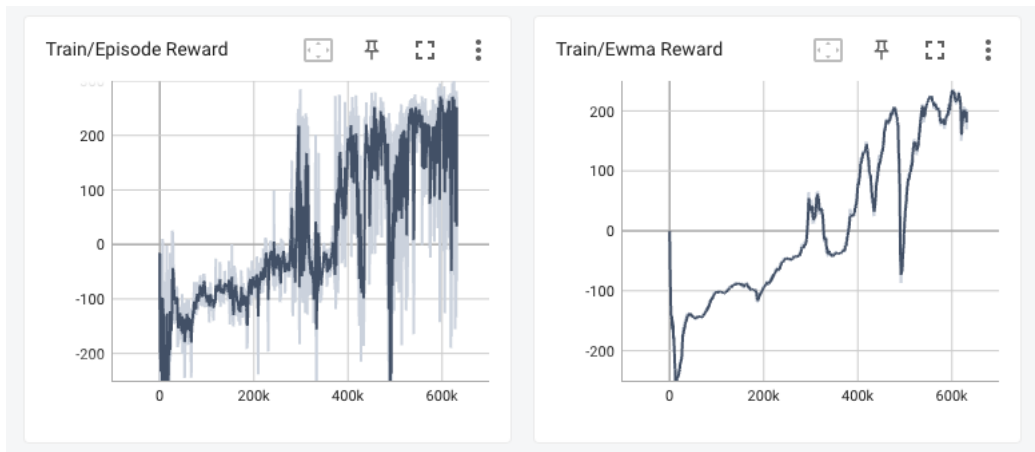
- Tensorboard



- Experimental Results of bonus parts (DDQN)
 - Testing Result

```
(dqn) minnie-mu@user:/eva_data6/minnie/NYCU/DLP/Lab6$ python ddqn.py --test_only
/home/minnie-mu/anaconda3/envs/dqn/lib/python3.8/site-packages/gym/logger.py:30: UserWarning:
  warnings.warn(colorize('%s: %s'%( 'WARN', msg % args), 'yellow'))
Start Testing
Episode: 0      \Length: 310      Total Reward: 222.20
Episode: 1      \Length: 310      Total Reward: 262.74
Episode: 2      \Length: 320      Total Reward: 259.48
Episode: 3      \Length: 385      Total Reward: 243.06
Episode: 4      \Length: 286      Total Reward: 296.61
Episode: 5      \Length: 292      Total Reward: 255.47
Episode: 6      \Length: 252      Total Reward: 178.17
Episode: 7      \Length: 269      Total Reward: 279.02
Episode: 8      \Length: 128      Total Reward: -88.27
Episode: 9      \Length: 549      Total Reward: 251.24
Average Reward 215.97105394196893
```

- Tensorboard



- Questions

- Describe your major implementation of both DQN and DDPG in detail. Your description should at least contain three parts:

1. Your implementation of Q network updating in DQN.

We use epsilon-greedy strategy to select action. If random number less than the epsilon, we will random sample an action, otherwise we will use behavior net to choose which action can get the maximum expected Q value.

```
def select_action(self, state, epsilon, action_space):
    if random.random() < epsilon:
        return action_space.sample()

    with torch.no_grad():
        state = torch.tensor(state, device=self.device).view(1, -1)
        outputs = self._behavior_net(state)
        _, best_action = torch.max(outputs, 1)
    return best_action.item()
```

We will sample random minibatch of transition $(\phi_j, a_j, r_j, \phi_{j+1})$ from replay memory. And compute target value y_j .

```
def _update_behavior_network(self, gamma):
    # sample a minibatch of transitions
    state, action, reward, next_state, done = self._memory.sample(
        self.batch_size, self.device)
    q_value = self._behavior_net(state).gather(1, action.long())
    with torch.no_grad():
        q_next = torch.max(self._target_net(next_state), 1)[0].view(-1, 1)
        q_target = reward + q_next * gamma * (1.0 - done)
    criterion = nn.MSELoss()
    loss = criterion(q_value, q_target)
    # optimize
    self._optimizer.zero_grad()
    loss.backward()
    nn.utils.clip_grad_norm_(self._behavior_net.parameters(), 5)
    self._optimizer.step()
```

Copy weight of behavior net to update target network every C steps.

```
def _update_target_network(self):
    '''update target network by copying from behavior network'''
    self._target_net.load_state_dict(self._behavior_net.state_dict())
```

2. Your implementation and the gradient of actor updating in DDPG.

We want to let behavior actor network predict an action which can get maximum q value by critic network. Use actor network to generate actions, and then get q value by critic network. Compute the mean q value and back-propagation to update actor network. Therefore we define actor loss as follows.

- s : current state
- Q : behavior critic network.
- u : behavior actor network

$$\text{Actor Loss} = -Q(s, u(s))$$

Update the actor policy using the sampled gradient:

$$\nabla_{\theta^\mu} \mu|_{s_i} \approx \frac{1}{N} \sum \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

3. Your implementation and the gradient of critic updating in DDPG.

First sample random minibatch from replay memory, and use actions of batch to calculate corresponding q value by critic network. We also need to calculate target value. First, get action of next state by target actor network, and then generate q value of next state by target critic network. Multiply discount factor and plus reward, we can get the target q value.

Set $y - i = \gamma_i + \gamma Q'(s_{t+1}, \mu'(s_{t+1}|\theta^{\mu'})|\theta^{Q'})$ Update critic by minimizing the loss; $L = \frac{1}{N} \sum (y_i - Q(s_i, a_i|\theta^Q))^2$

4. Update behavior network: We will sample random minibatch (s_i, a_i, r_i, s_{i+1}) from replay memory. And compute target value y_i , the formulation is as follows.

- Q : behavior critic network
- u : behavior actor network
- \hat{Q} : target critic network
- \hat{u} : target actor network
- γ : discount factor
- $y_i = r_i + \gamma * \hat{Q}(s_{i+1} + \hat{u}(s_{i+1}))$

And use mean square error as loss function to update behavior critic network.

$$\text{MSE Loss} = \frac{1}{N} \sum (y_i - Q(s_i, a_i))^2$$

And update behavior actor network by following formulation.

$$\text{Actor Loss} = -\frac{1}{N} \sum (Q(s_i), u(s_i))$$

```
def _update_behavior_network(self, gamma):
    actor_net, critic_net, target_actor_net, target_critic_net = self._actor_net, self._critic_net, self._target_actor_net, self._target_critic_net
    actor_opt, critic_opt = self._actor_opt, self._critic_opt

    # sample a minibatch of transitions
    state, action, reward, next_state, done = self._memory.sample(
        self.batch_size, self.device)

    ## update critic ##
    # critic loss
    q_value = critic_net(state, action)
    with torch.no_grad():
        a_next = target_actor_net(next_state)
        q_next = target_critic_net(next_state, a_next)
        q_target = reward + gamma * q_next * (1 - done)
    criterion = nn.MSELoss()
    critic_loss = criterion(q_value, q_target)

    # optimize critic
    actor_net.zero_grad()
    critic_net.zero_grad()
    critic_loss.backward()
    critic_opt.step()

    ## update actor ##
    # actor loss
    action = actor_net(state)
    actor_loss = -critic_net(state, action).mean()

    # optimize actor
    actor_net.zero_grad()
    critic_net.zero_grad()
    actor_loss.backward()
    actor_opt.step()
```

- Explain effects of the discount factor.

By the Q learning algorithm, γ is the discount factor. the smaller discount factor is, the agent will focus on current reward. The larger discount factor is, the agent will focus on future reward.

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_{a'} Q(S', a') - Q(S, A))$$

- Explain benefits of epsilon-greedy in comparison to greedy action selection.

The model choose the action which is the best for current state, but the result is not necessarily the best for the future. So, to let model explore other action, we use epsilon-greedy to set a probability ϵ which can let model to randomly choose other actions.

- Explain the necessity of the target network.

If we don't have target network, the loss is calculated with behavior q value and behavior target q value. This may cause difficulty for training. So, to avoid this situation, we need target network and copy weights to target network from behavior network every constant steps.

- Describe the tricks you used in Breakout and their effects, and how they differ from those used in LunarLander.

The difference between Breakout and LunarLander is the input, the input of Breakout include 4 frames and the LunarLander has just one frame.